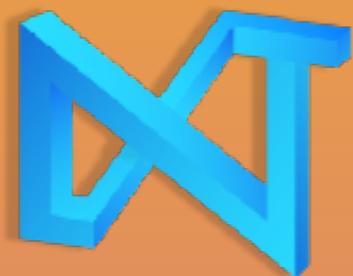


Delve Into Art and Creativity with Python

Aparna Krishnakumar

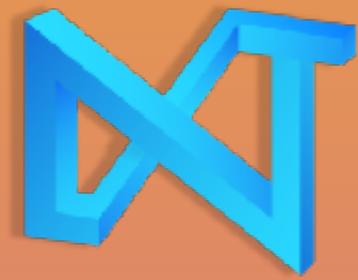


Next Tech Lab



SRM
INSTITUTE OF SCIENCE & TECHNOLOGY
—Deemed to be University—

SRM University



Next Tech Lab



SRM
INSTITUTE OF SCIENCE & TECHNOLOGY
—Deemed to be University—

SRM University

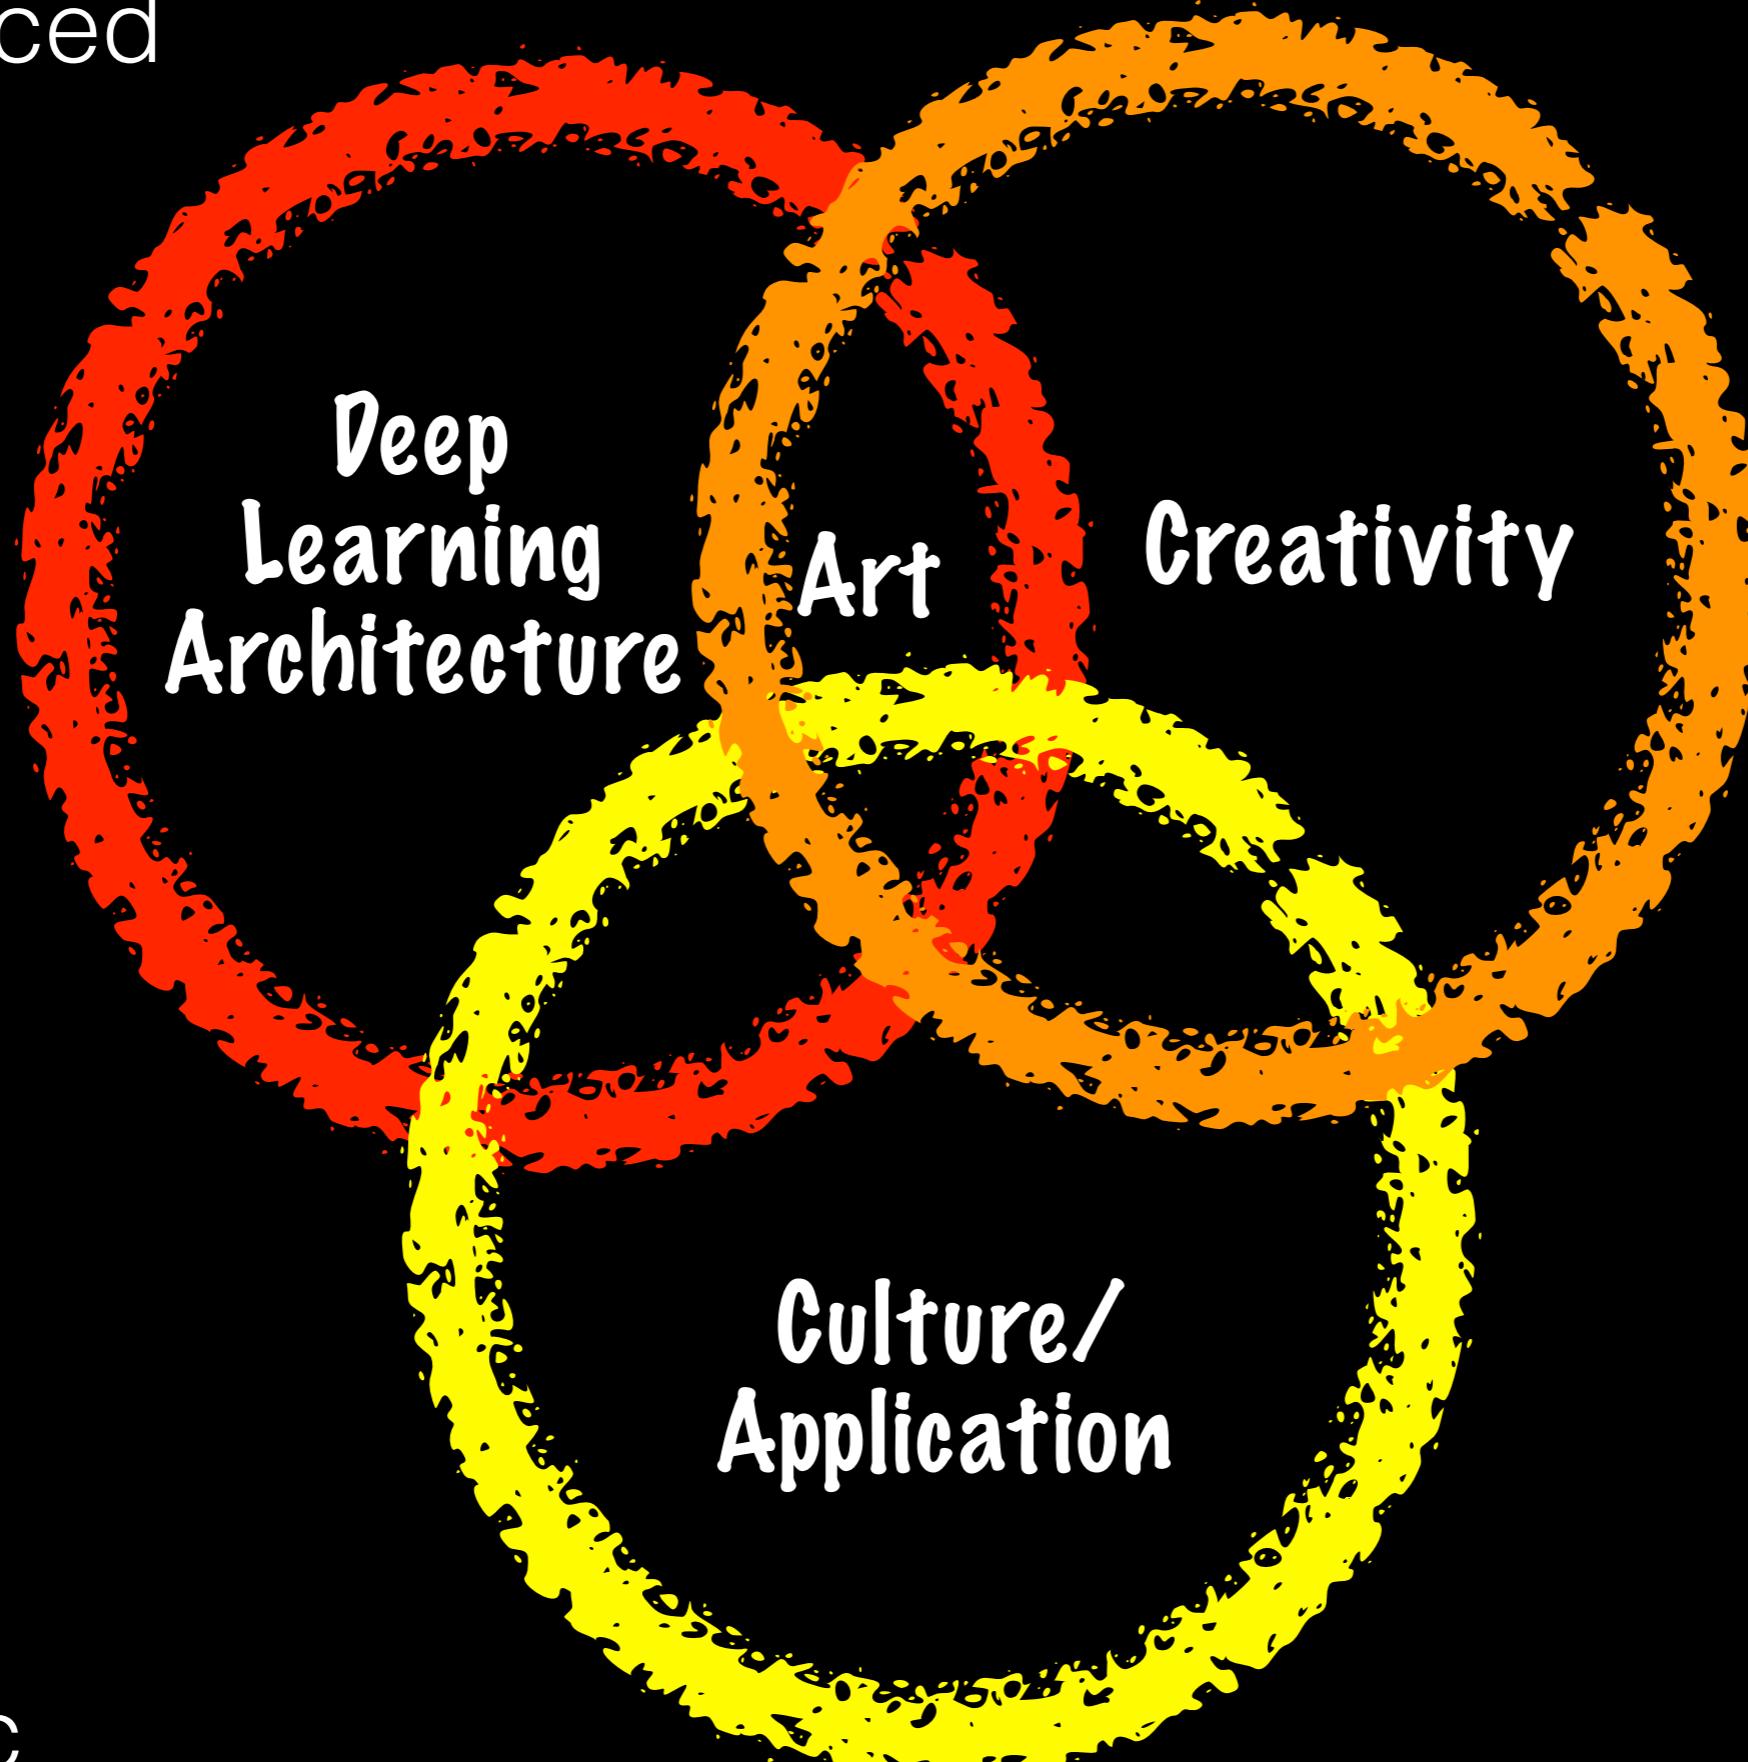
Deep Learning

what has been done and what can be done



Advanced

Culture/
Application



Basic

'ART'

Libraries

Why Python?

- Easy implementation
- Powerful, Image processing libraries
- comprehensive material
- Open Source
- AI Frameworks- TensorFlow, Keras, Pytorch
- Community
- Research Oriented

Artificial Intelligence-Deep Learning

Data

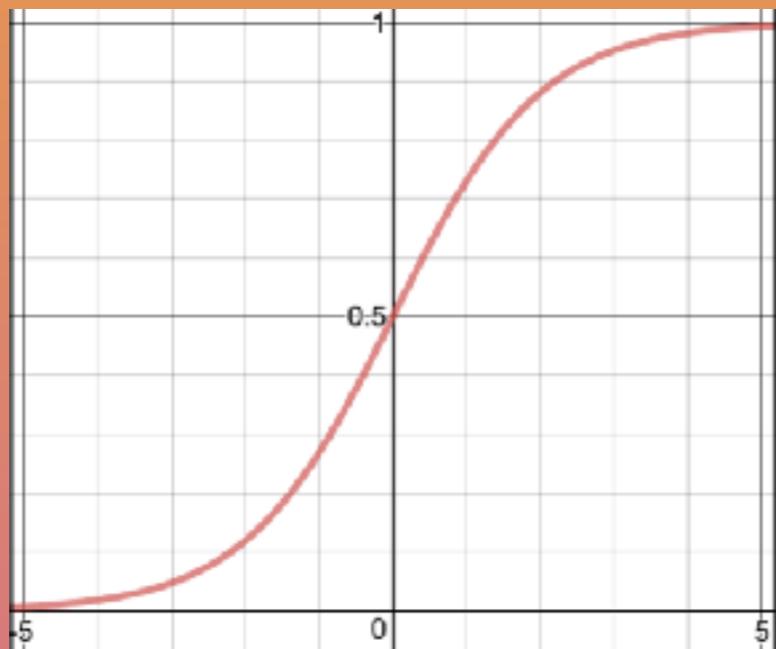
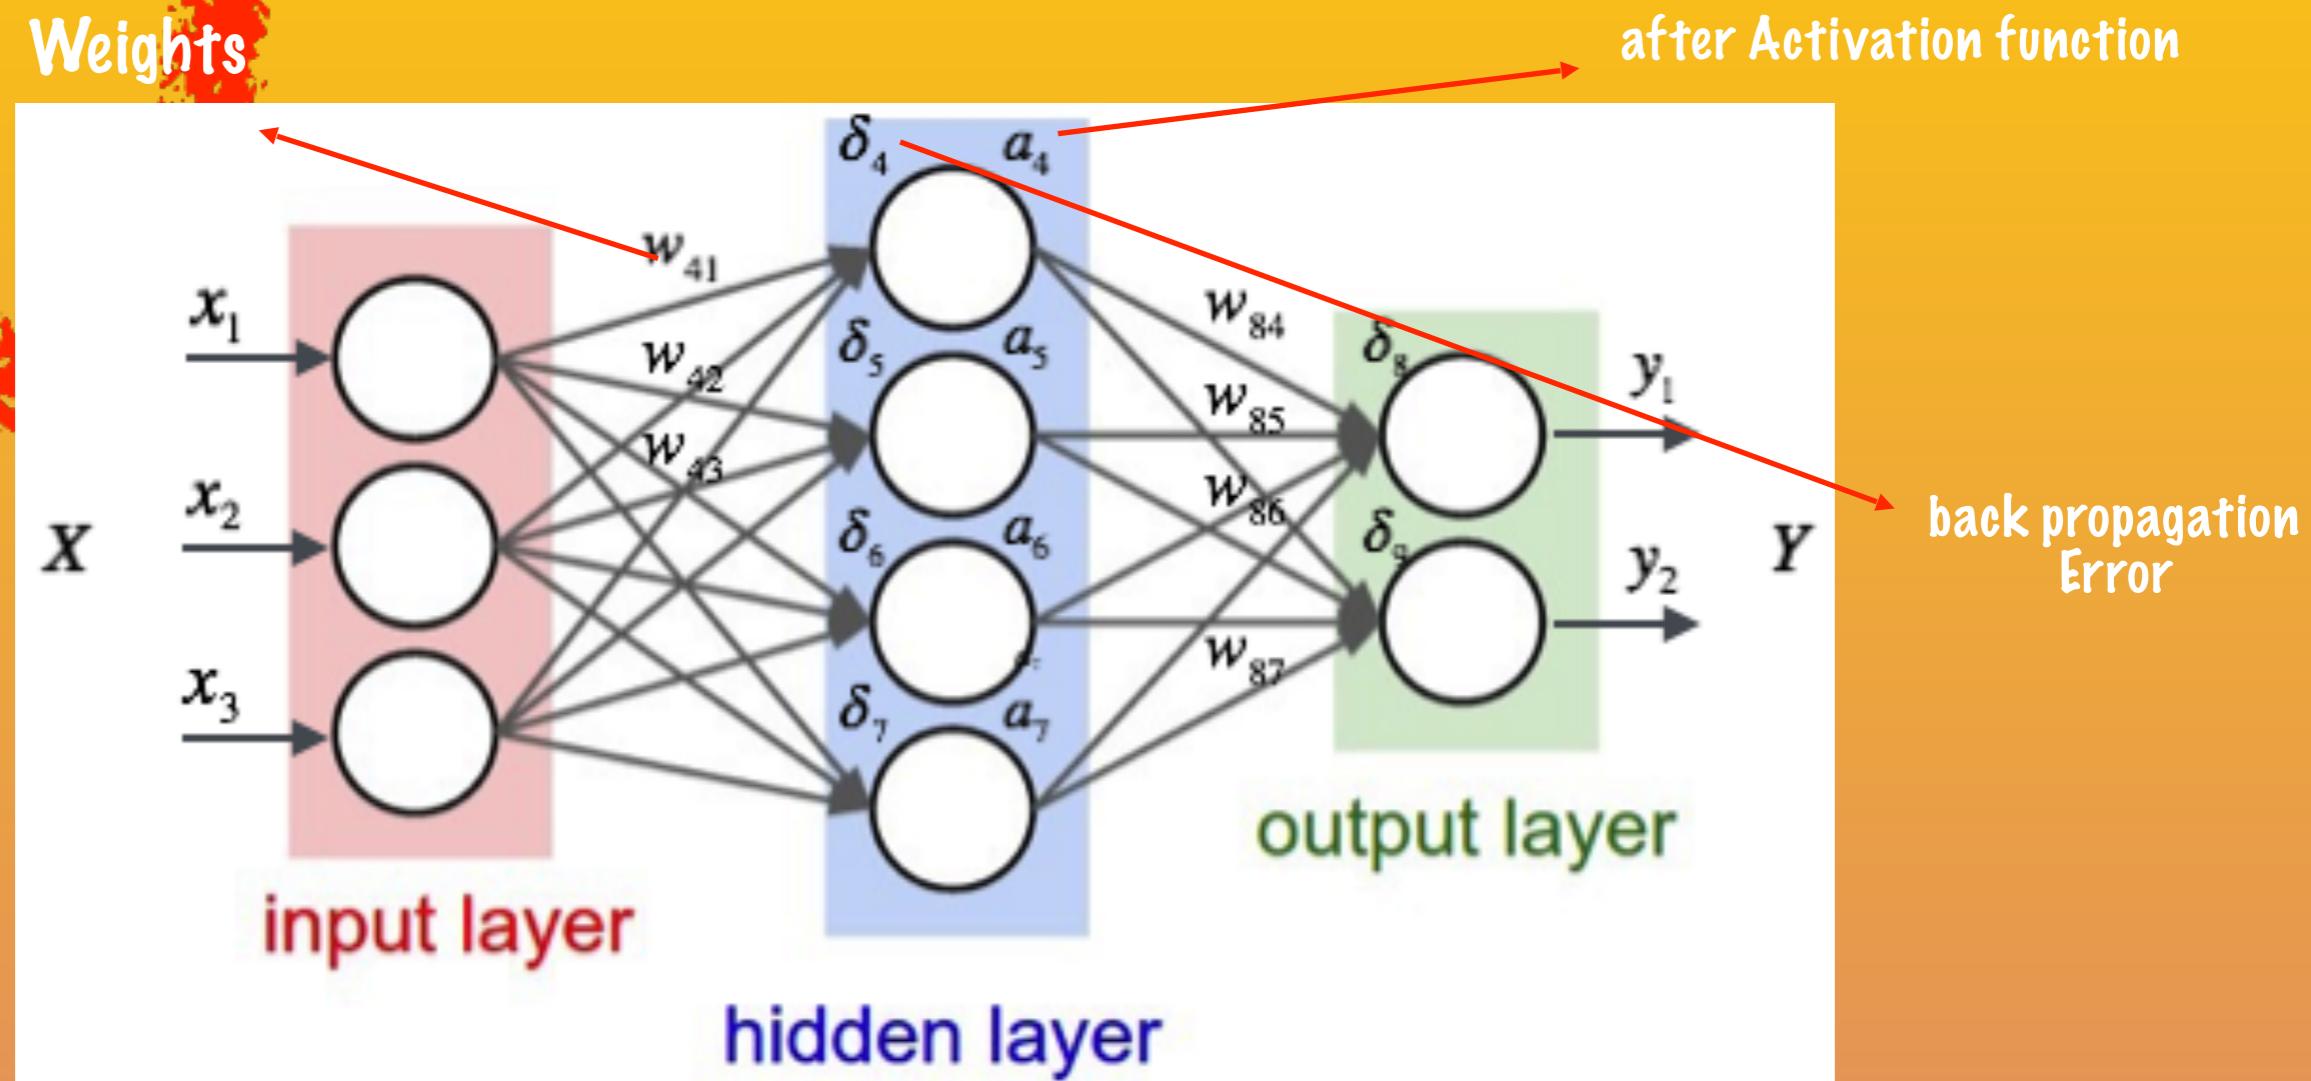
m -by- n matrix

$a_{i,j}$ n columns j changes

m rows	i changes	
$a_{1,1}$	$a_{1,2}$	$a_{1,3}$. . .
$a_{2,1}$	$a_{2,2}$	$a_{2,3}$. . .
$a_{3,1}$	$a_{3,2}$	$a_{3,3}$. . .
:	:	:
:	:	:

Training examples X features

Neural Networks



$$a = g(w^*x + b)$$

Sigmoid function

$$S(z) = \frac{1}{1 + e^{-z}}$$

cost function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))]$$

```

def sigmoid(z):
    """
    the sigmoid function is used as an activation function to convert linear outputs
    to non linear outputs such that probability is outputed in between 0 and 1
    """

    return 1/(1+np.exp(-z))

```

```

def forward_propagate(X,theta1,theta2):
    """
    This function propagates through the network computing the output of
    every layer of the neural network with the given inputs and weights and
    computes the final output of the xor gate
    """

    #first, a column of biases is added to the input of the first layer
    a1=np.c_[np.ones(X.shape[0]),X]#mxn+1
    #the weights of the first layer are multiplied by the input of the first layer
    z1=a1.dot(theta1)#mxhidden
    #the input of the second layer is the output of the first layer, passed through the activation
    #function and column of biases is added
    a2=np.c_[np.ones(X.shape[0]),sigmoid(z1)]#mxhidden+1
    #the input of the second layer is multiplied by the weights
    z2=a2.dot(theta2)
    #the output is passed through the activation function to obtain the final probability
    h3=sigmoid(z2)
    #print(h3.shape)
    return a1,z1,a2,z2,h3

```

```

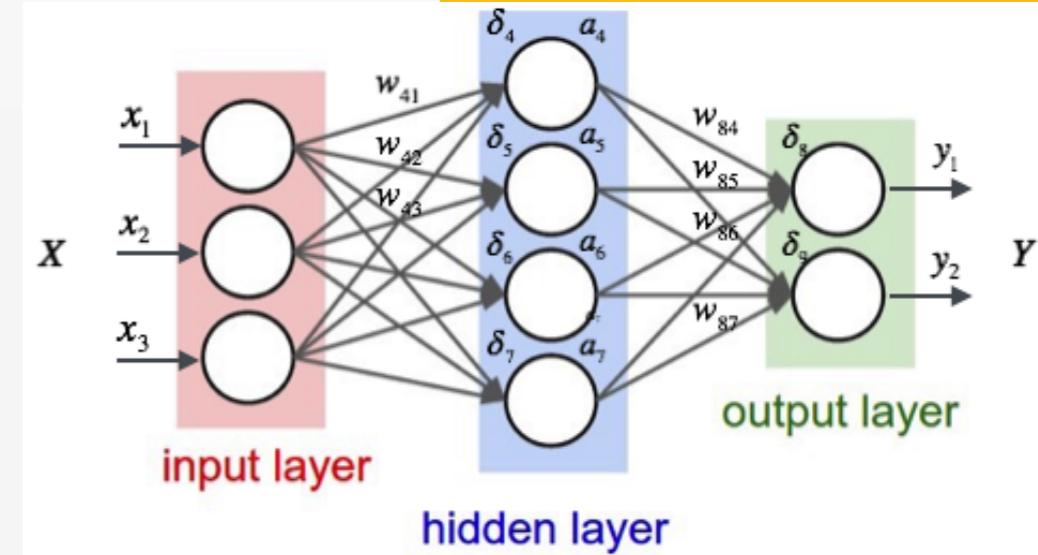
for i in range(1000):
    a1, z1, a2, z2, hyp = forward_propagate(X, theta1, theta2)# for every iteration, forward propagation is carried out
    del_2=Y-hyp#the error of the final layer is calculated- the difference between the predicted
    and actual output
    #the error of the previous layer is found by computing the dot product of the error of the previous layer and the weights of the second layer,without the column for biases.
    #this matrix is made to undergo element-wise multiplication with the output of the first layer
    #(taking into account the activation function)
    del_1=del_2.dot(theta2[1:,:].T)

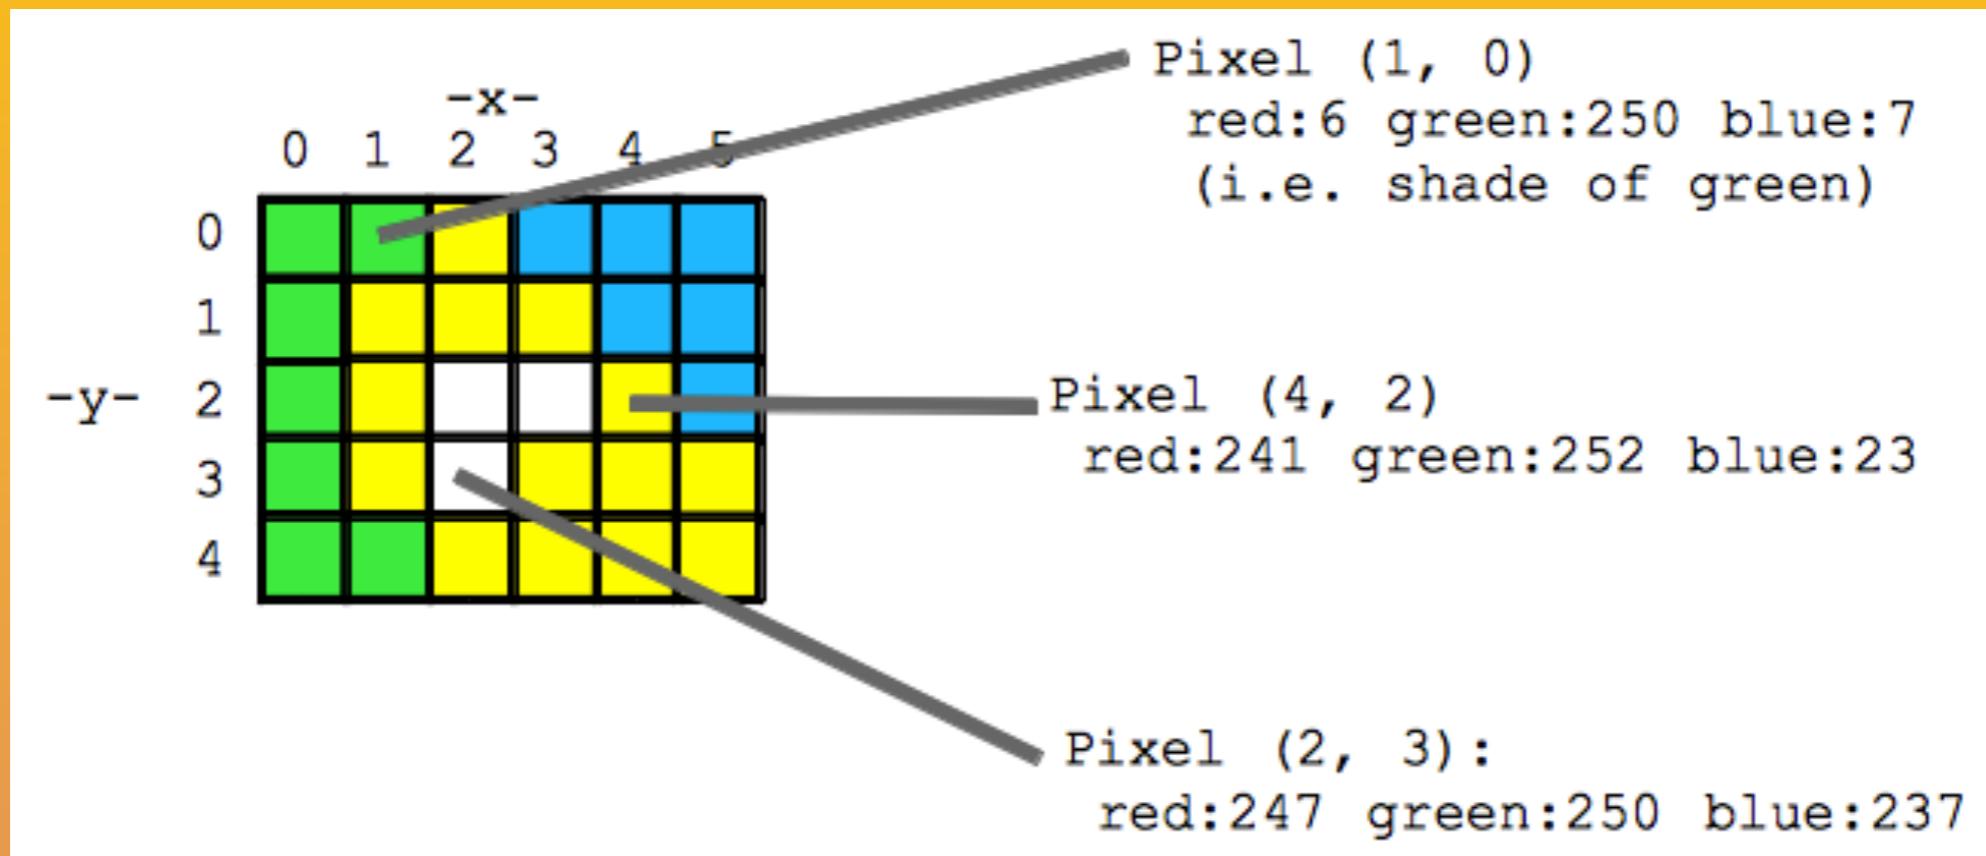
    #the error of the second layer is multiplied element wise by the sigmoid gradient of the output
    #of the second layer
    delta2=del_2

    #the error of the first layer is multiplied element wise by the sigmoid gradient of the output
    #of the first layer
    delta1=del_1*sigmoid_grad(z1)

    #the parameters are updated using gradient descent
    theta2+=l_r*a2.T.dot(delta2)
    theta1+=l_r*a1.T.dot(delta1)

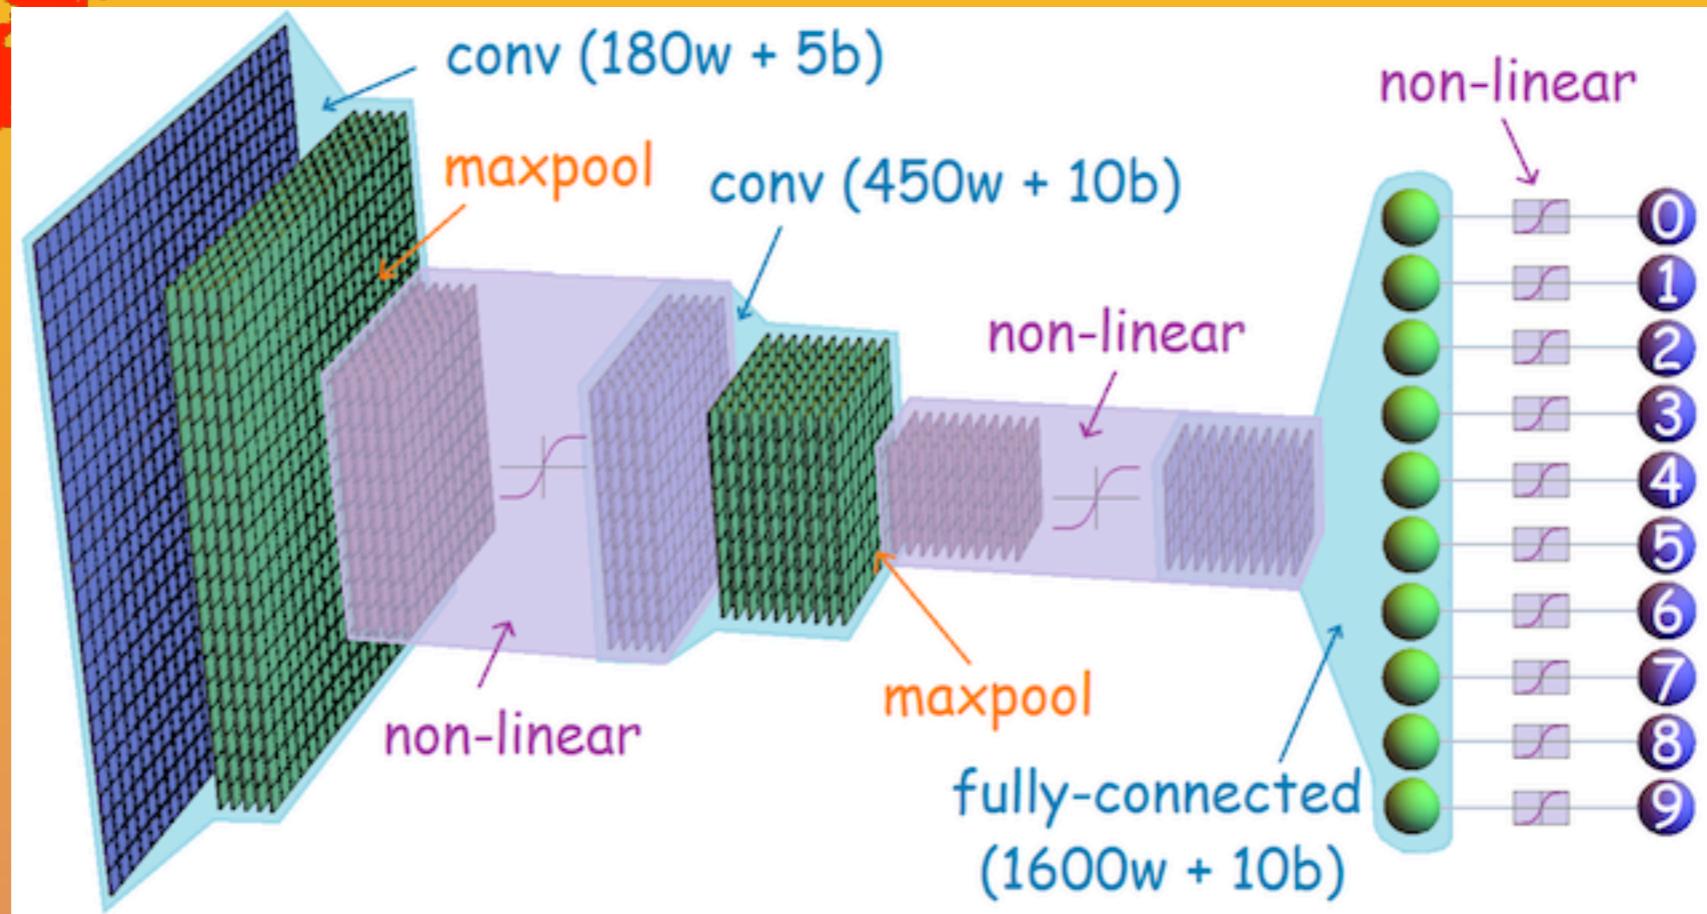
```





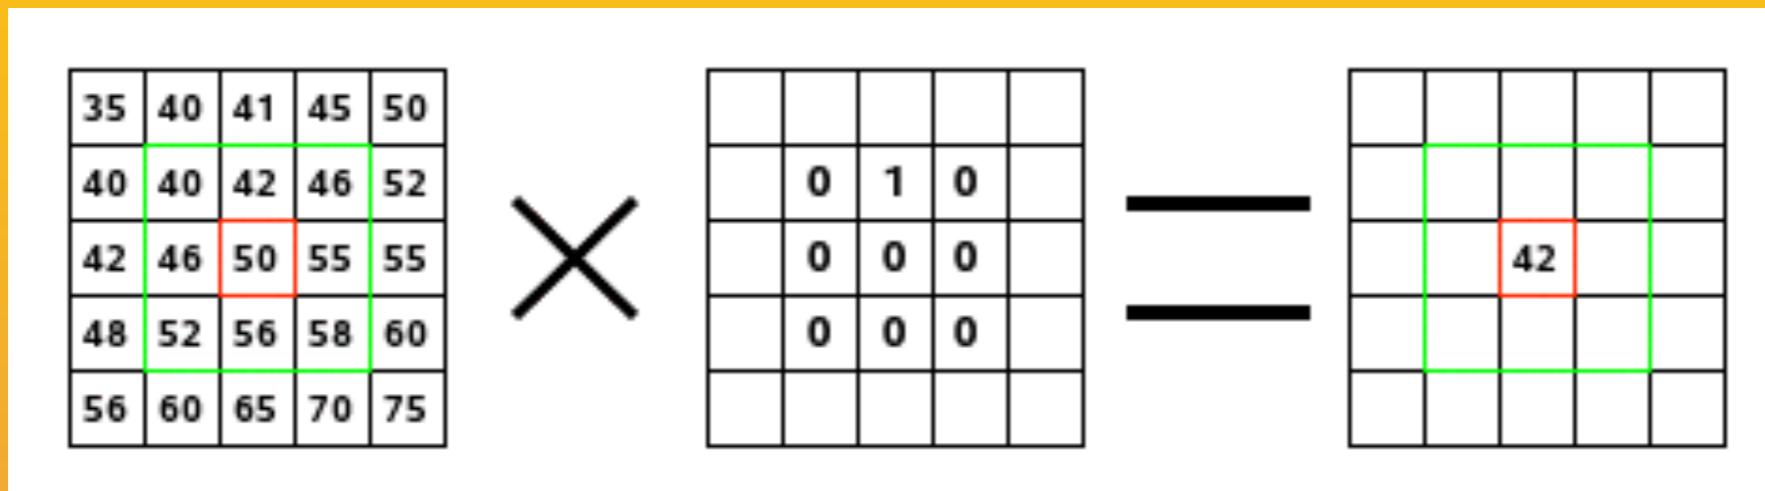
pixel values :0-255
Height X width X channel

Convolutional Neural Networks

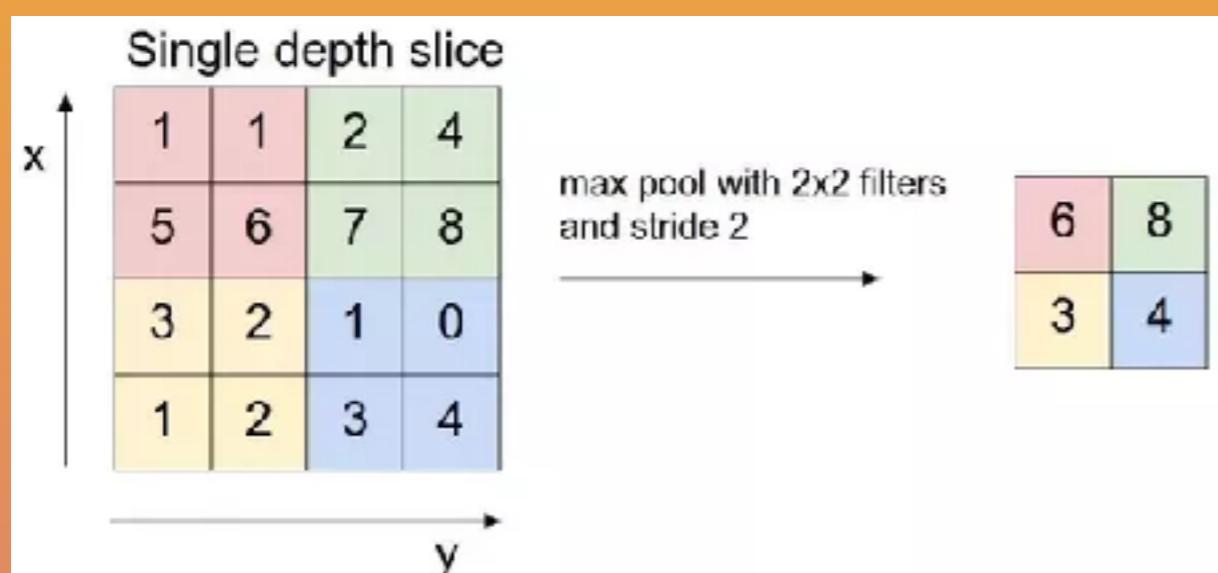


Training examples X features

1. Convolution



2. Max Pooling



3. RELU

4. Fully connected Layer

```

def forward_propagation(X, parameters):
    W1 = parameters['W1']
    W2 = parameters['W2']

    Z1 = tf.nn.conv2d(X, W1, strides=[1, 1, 1, 1], padding='SAME')
    # RELU
    A1 = tf.nn.relu(Z1)
    # MAXPOOL
    P1 = tf.nn.max_pool(A1, ksize=[1, 8, 8, 1], strides=[1, 8, 8, 1], padding='SAME')
    # CONV2D
    Z2 = tf.nn.conv2d(P1, W2, strides=[1, 1, 1, 1], padding='SAME')
    # RELU
    A2 = tf.nn.relu(Z2)
    # MAXPOOL
    P2 = tf.nn.max_pool(A2, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1], padding='SAME')
    # FLATTEN
    P = tf.contrib.layers.flatten(P2)

    Z3 = tf.contrib.layers.fully_connected(P, 6, activation_fn=None)

    return Z3

```

```

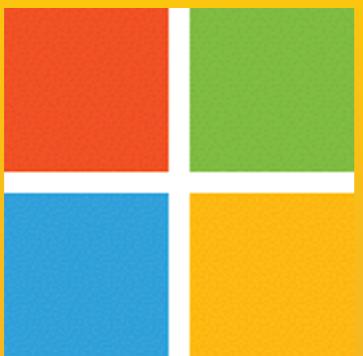
def compute_cost(Z3, Y):
    cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=Z3, labels=Y))

    return cost

```

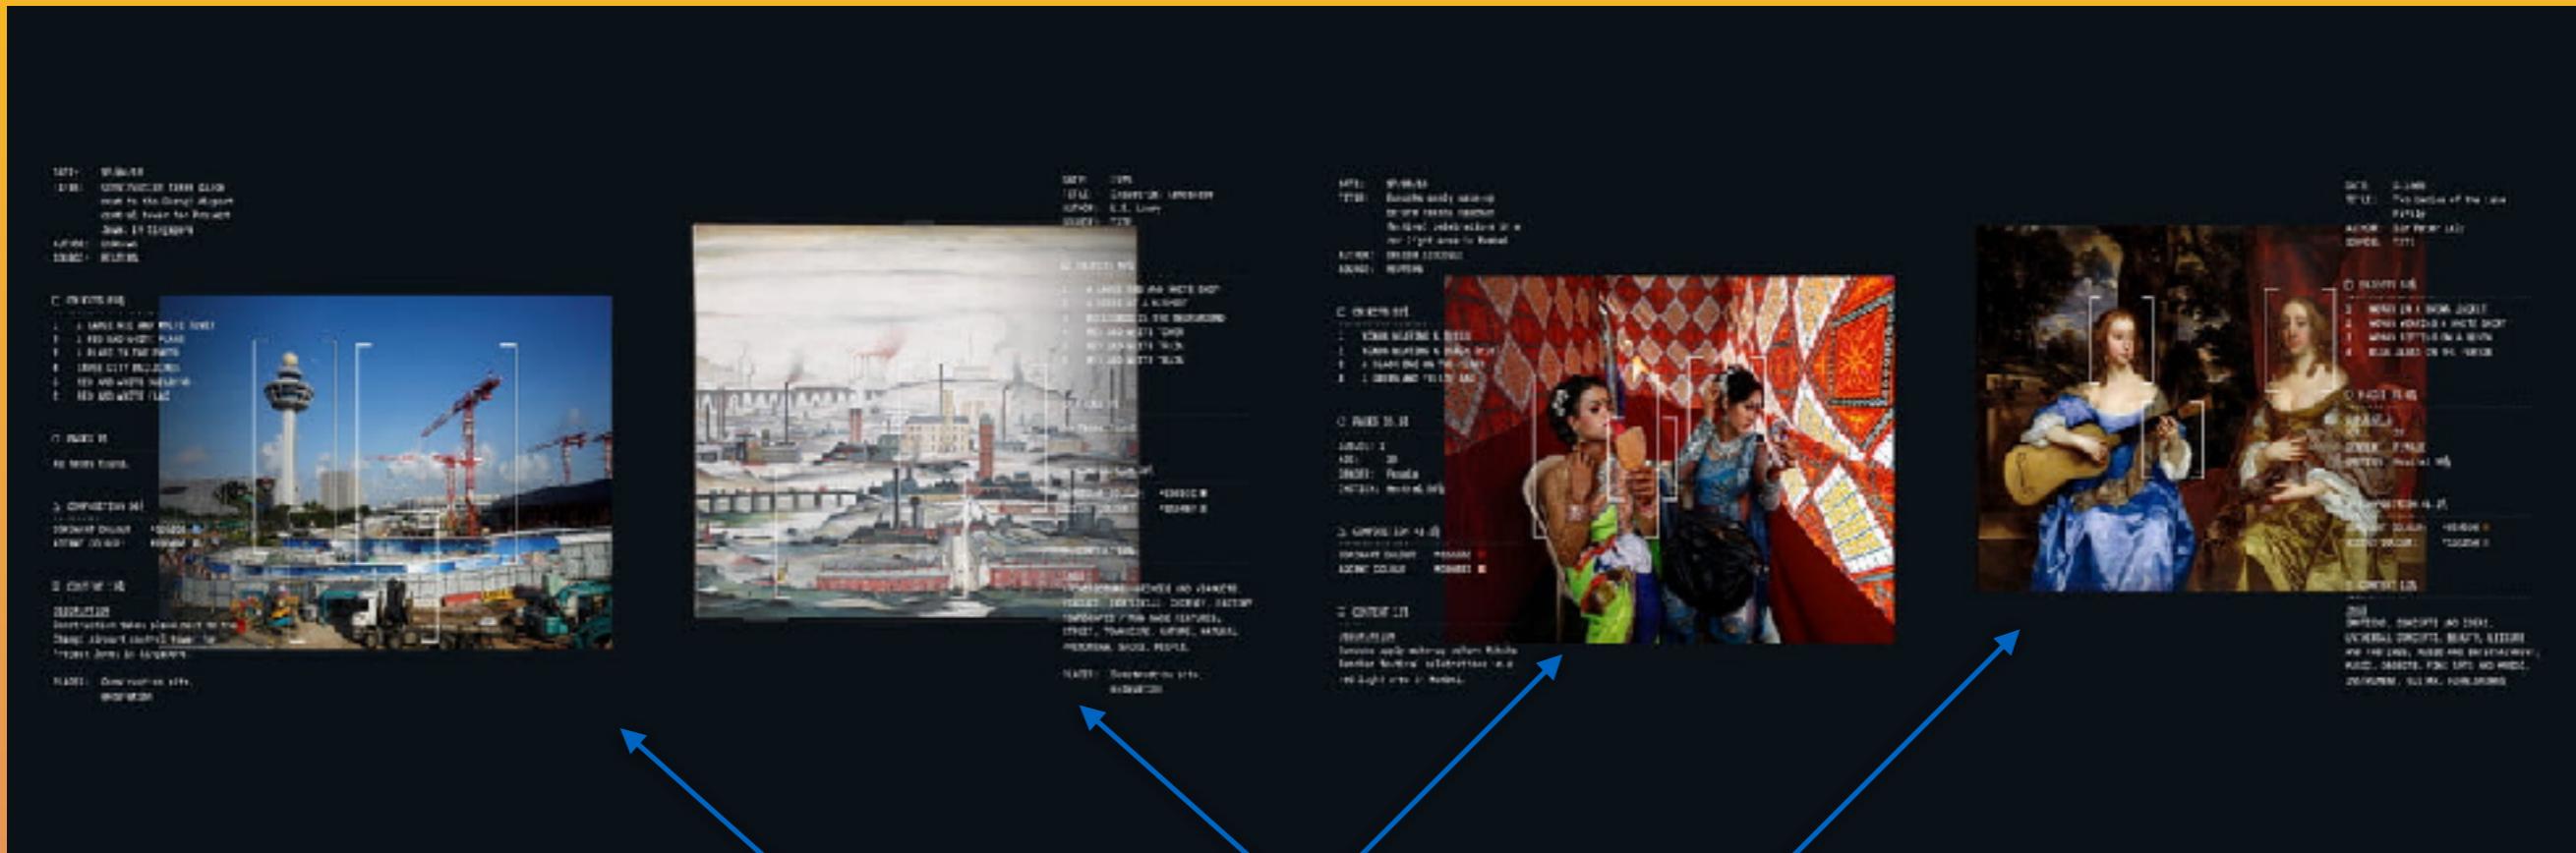
```
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

Convl Conv2 Conv3 Conv4
 → more detail



Tate Britain AI exhibit

Recognition 2016



Reuters

Tate

Style Transfer

Leon A. Gatys, Alexander S. Ecker, Matthias Bethge



Conv 4

Alpa, beta

Max Pooling

Gram Matrix

Image1

Content

$$\mathcal{L}_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2 .$$

Style

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l .$$

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

$$\mathcal{L}_{style}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l$$

Experiment

Style on style
genetic algorithms
Conv 2?
other types of pooling?

$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

Noise

Using weights that
minimise the cost function

Output

Style Transfer + Culture

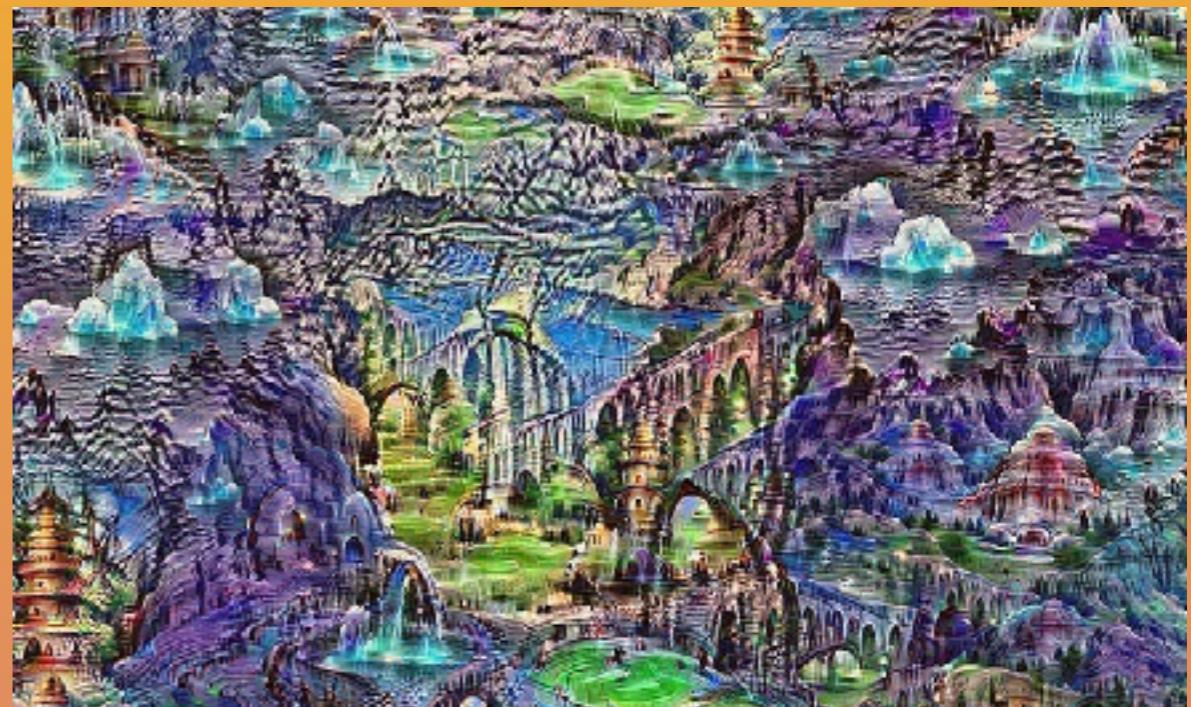


Deep Greetings

In masks and gowns we haunt the street,
and knock on doors for trick or treat.
tonight we are the king and queen,
for oh tonight it's halloween!



Deep Dream



Experiment

Hybrid organisms
Fuse cultures
Scans/ medical Imaginary

Iterate over Outputs:

Features

Snapshot of
inner workings

feed as input

Foster The People

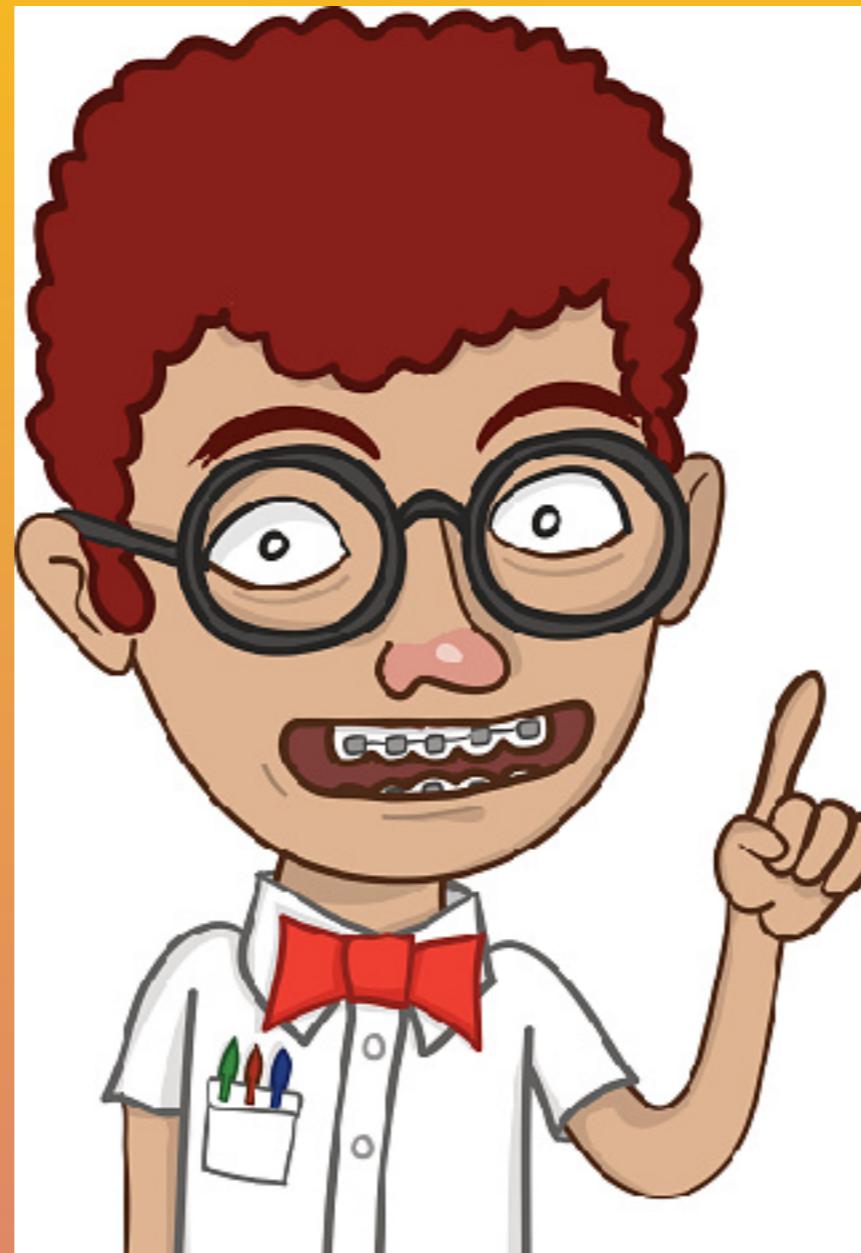
Doing it for the Money



GANS



Discriminator



Generator



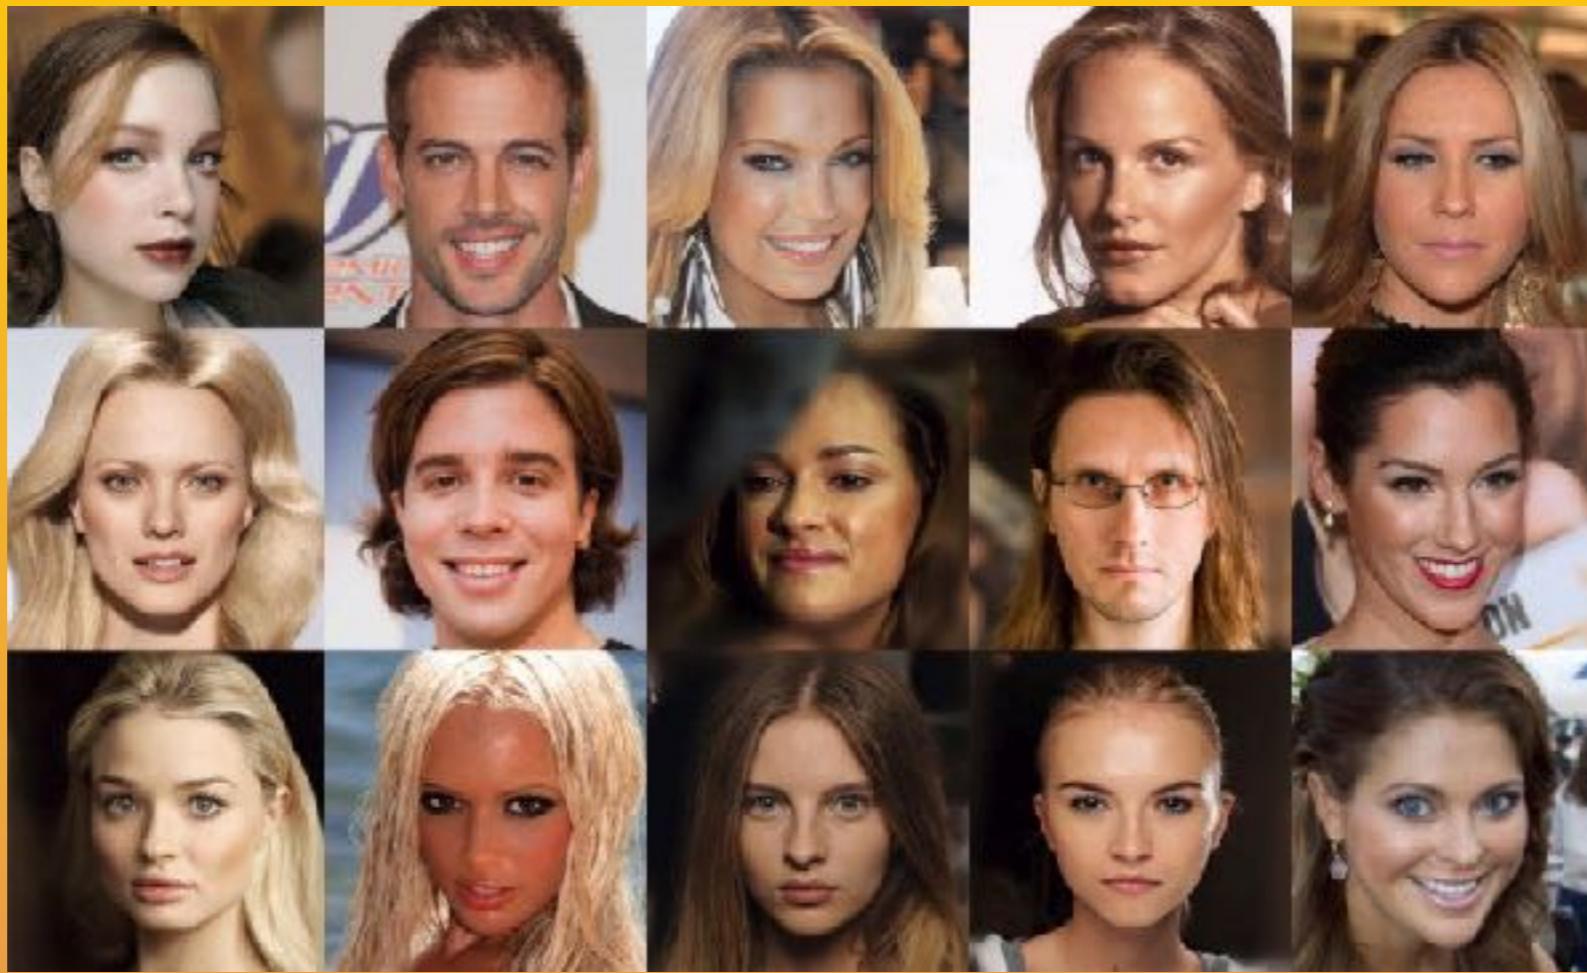
Training Data



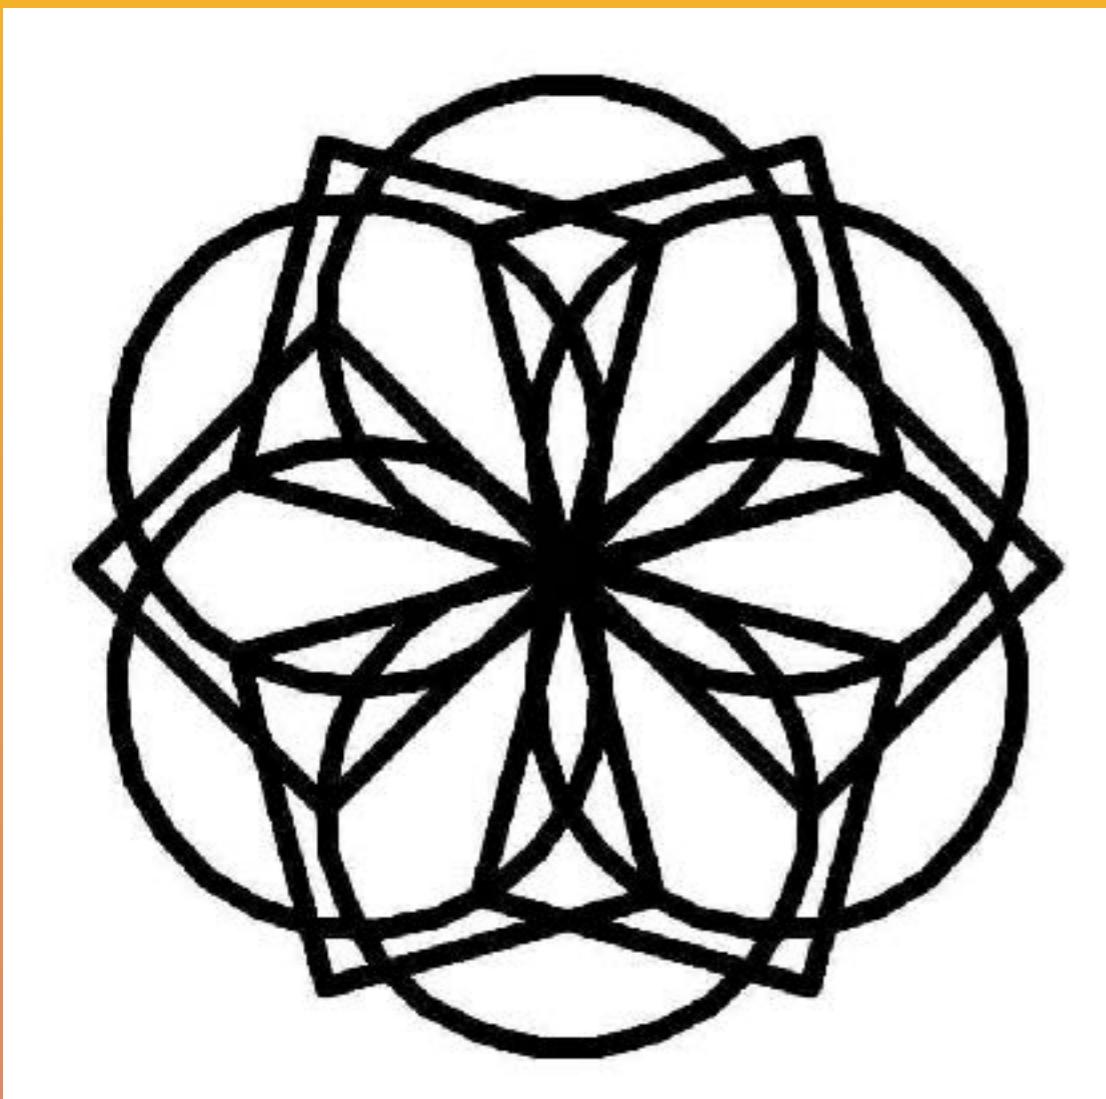
```
def generator_fn(noise, weight_decay=2.5e-5, is_training=True):
    with framework.arg_scope(
        [layers.fully_connected, layers.conv2d_transpose],
        activation_fn=tf.nn.relu, normalizer_fn=layers.batch_norm,
        weights_regularizer=layers.l2_regularizer(weight_decay)), \
        framework.arg_scope([layers.batch_norm], is_training=is_training, zero_debias_moving_mean=True):
        net = layers.fully_connected(noise, 1024)
        net = layers.fully_connected(net, 800 * 800 * 1)
        net = tf.reshape(net, [-1, 800, 800, 1])
        net = layers.conv2d_transpose(net, 64, [4, 4], stride=2)
        net = layers.conv2d_transpose(net, 32, [4, 4], stride=2)
        # Make sure that generator output is in the same range as `inputs`
        # ie [-1, 1].
        net = layers.conv2d(net, 1, 4, normalizer_fn=None, activation_fn=tf.tanh)

    return net


def discriminator_fn(img, unused_conditioning, weight_decay=2.5e-5,
                     is_training=True):
    with framework.arg_scope(
        [layers.conv2d, layers.fully_connected],
        activation_fn=leaky_relu, normalizer_fn=None,
        weights_regularizer=layers.l2_regularizer(weight_decay),
        biases_regularizer=layers.l2_regularizer(weight_decay)):
        net = layers.conv2d(img, 64, [4, 4], stride=2)
        net = layers.conv2d(net, 128, [4, 4], stride=2)
        net = layers.flatten(net)
        with framework.arg_scope([layers.batch_norm], is_training=is_training):
            net = layers.fully_connected(net, 1024, normalizer_fn=layers.batch_norm)
    return layers.linear(net, 1)
```



Deep Rangoli



Coded Couture



- Abstract
- Style Transfer
- Creative Coding
- GANS
- Augmented Intelligence



Why should this be explored further?

- Evolution, expression
- Innovation
- applications
- Architectures
- logic
- socio-economic applications

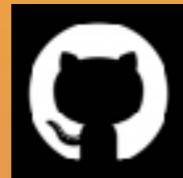
“ there's more than one way to do it”



@Aparnaakk



aparnakrishnakumar.com



@Aparnaakk